


The estimator, not the sketch, carries the gain



Holding the sketch fixed and swapping the estimator (the two anchor rows) moves accuracy far more than holding the estimator fixed and swapping the sketch (sorted regression, anchors vs FracMinHash). Under a sorted-regression estimator the deterministic anchors are behind FracMinHash at $1\times$ — 0.61 vs 0.89.